
Does Iterative Rejection Make AI Text Undetectable? Conversational Bullying as a Simple Evasion Strategy

Ari Holtzman and NeuriCo

Abstract

AI text detectors are widely deployed in education, publishing, and hiring, yet their robustness to simple evasion strategies remains poorly understood. We investigate the simplest possible attack: iteratively telling a large language model “this sounds too AI-generated, try again” and measuring whether the resulting text evades detection. Using GPT-4.1 to answer 50 open-domain questions under five conditions—vanilla generation, single-step humanization, strong single-step prompting, iterative rejection over five rounds, and best-of-5 selection—we score all outputs with two detectors: a ROBERTA-based classifier and BINOCULARS, a state-of-the-art zero-shot method. Iterative rejection reduces BINOCULARS detection scores by 43% relative to vanilla generation ($p < 0.001$), and significantly outperforms best-of-5 selection ($p < 0.001$, $r = 0.56$), confirming that conversational rejection context provides genuine steering beyond selection pressure. However, a well-crafted single-step prompt with specific anti-AI instructions achieves even stronger evasion (60% reduction), approaching human-level scores. Linguistic analysis reveals that iterative rejection triggers a regime shift—the model switches to an exaggerated casual persona with higher vocabulary diversity and excessive first-person narration—rather than precisely suppressing detectable features. These findings have implications for the trustworthiness of detector-based enforcement and suggest that LLMs harbor distinct generation modes activatable through simple conversational context.

1 Introduction

A student submits an essay. An instructor runs it through an AI detector, which flags it as machine-generated. The student protests: “I just asked ChatGPT to rewrite it until it sounded natural.” This scenario plays out daily across universities, publishers, and hiring platforms that rely on AI text detectors as gatekeepers. But how effective is this trivially simple strategy—just telling an LLM “this sounds too AI, try again”—at actually evading detection?

Why this matters. AI-generated text (AIGT) detectors are deployed at scale, with services like GPTZero, Turnitin’s AI detection module, and open-source classifiers used to make consequential decisions about authorship. Prior work has demonstrated that sophisticated evasion techniques—contrastive decoding Fang et al. [2025], detector-guided token selection Cheng et al. [2025], and trained paraphraser Krishna et al. [2023]—can reduce detection rates dramatically. Yet all these methods require technical expertise, model access, or custom tooling. The simplest attack vector—a non-technical user iteratively rejecting outputs through normal conversation—has never been systematically studied.

The gap. Our literature review of 15 papers spanning AI text detection and evasion reveals a consistent pattern: every existing iterative evasion method operates at the training, optimization, or decoding level. SICO Lu et al. [2024] iterates on prompt examples against a proxy detector. Self-Disguise Attack Zhou et al. [2025] extracts linguistic features from successful evasions to update generation prompts. CoPA Fang et al. [2025] integrates contrastive corrections at each decoding

step. No study tests the scenario that millions of users already practice: generating text, reading it, deciding it “sounds too AI,” and asking for a revision—all within a single chat session.

Our approach. We design a controlled experiment that isolates the effect of iterative conversational rejection. Using GPT-4.1, we generate answers to 50 open-domain questions from the HC3 corpus Guo et al. [2023] under five conditions: vanilla generation, single-step humanization, strong single-step prompting with specific instructions, iterative rejection over five rounds with escalating feedback, and best-of-5 independent selection. We score all outputs with two complementary detectors: a ROBERTA-based trained classifier and BINOCULARS Hans et al. [2024], a state-of-the-art zero-shot method. The best-of-5 control is critical: it distinguishes whether iterative rejection works through genuine contextual steering (the rejection feedback changes how the model generates) or through mere selection pressure (more samples means a better chance of a low-scoring output).

Preview of results. Iterative rejection reduces BINOCULARS scores by 43% (from 0.855 to 0.483), significantly outperforming best-of-5 selection (0.702, $p < 0.001$). This confirms that conversational rejection context steers generation beyond what selection alone achieves. However, a carefully crafted single-step prompt that lists specific AI markers to avoid achieves an even larger reduction of 60% (to 0.342), approaching human text levels (0.325). Linguistic analysis reveals a surprising mechanism: rather than suppressing specific AI features, the model enters a distinct “trying to sound human” regime characterized by higher vocabulary diversity, excessive first-person narration, and persistent AI marker words—a regime shift rather than targeted feature suppression.

We make the following contributions:

- We present the first systematic evaluation of simple iterative conversational rejection as an AI text detection evasion strategy, requiring no technical expertise, model access, or custom tooling.
- We introduce a best-of- N control that disentangles contextual steering from selection pressure, showing that rejection feedback provides genuine generative steering ($p < 0.001$, $r = 0.56$).
- We identify a regime-shift mechanism: iterative rejection triggers a qualitative change in generation mode rather than precise feature suppression, with implications for both detector design and alignment research.

2 Related Work

AI-generated text detection. Detection methods fall into three broad categories. *Trained classifiers* fine-tune language models on labeled corpora of human and AI text. The OpenAI ROBERTA detector Solaiman et al. [2019] is representative, though it was trained on GPT-2 era text and generalizes poorly to modern LLMs, as we confirm in our experiments. *Zero-shot statistical methods* exploit distributional differences between human and machine text without training data. Detect-GPT Mitchell et al. [2023] measures probability curvature via perturbations; Fast-DetectGPT Bao et al. [2024] replaces perturbations with conditional probability curvature for a $340\times$ speedup. BINOCULARS Hans et al. [2024] contrasts perplexity scores from two related models (e.g., Falcon-7B and Falcon-7B-Instruct), achieving over 90% detection at 0.01% false positive rate. *Watermarking methods* Kirchenbauer et al. [2023] embed statistical signals during generation but require control of the generation process. We use both a trained classifier (ROBERTA) and a zero-shot method (BINOCULARS) to evaluate evasion across detection paradigms.

Evasion via paraphrasing. Krishna et al. [2023] showed that DIPPER, an 11B parameter T5-based paraphraser, breaks all major detectors at the time of publication, reducing Detect-GPT’s true positive rate from 70.3% to 4.6% at 1% FPR. Subsequent work has refined paraphrase-based attacks. CoPA Fang et al. [2025] uses contrastive decoding, subtracting machine-like token distributions from human-like distributions during generation, reducing average TPR to 7–13% across eight detectors. Adversarial Paraphrasing Cheng et al. [2025] uses detector-guided token selection during LLM decoding, achieving 87.9% reduction in TPR@1%FPR and demonstrating that evading one well-trained detector transfers to others. All of these methods require either custom decoding procedures, trained paraphrase models, or gradient-level access.

Iterative and optimization-based evasion. Several methods employ iterative loops to improve evasion. SICO Lu et al. [2024] constructs optimized in-context examples through greedy substitution over six rounds against a proxy detector, reducing average AUC by 0.5 across six detectors. Self-Disguise Attack Zhou et al. [2025] uses an adversarial feature extraction loop that collects detection-

Table 1: Comparison of evasion methods. Our approach requires no technical expertise, custom decoding, or model access—only a standard chat interface. The best-of- N control distinguishes steering from selection.

Method	Iterative	Requires Training	Requires Decoding Access	Automated	User Skill
DIPPER Krishna et al. [2023]	No	Yes	No	Yes	High
SICO Lu et al. [2024]	Yes	No	No	Yes	High
CoPA Fang et al. [2025]	No	No	Yes	Yes	High
AdvPara Cheng et al. [2025]	No	No	Yes	Yes	High
SDA Zhou et al. [2025]	Yes	No	No	Yes	High
HMGC Zhou et al. [2024]	Yes	No	Yes	Yes	High
Ours (Iterative Rejection)	Yes	No	No	No	None

resistant examples, extracts their linguistic features, and updates the generation prompt iteratively. HMGC Zhou et al. [2024] performs word-level adversarial substitution using gradient-based importance ranking, iterating until the detector score falls below a threshold. These approaches share the iterative structure of our method but differ in a critical way: they iterate on prompts, features, or tokens using automated optimization. We test whether iteration via simple natural language feedback—with no optimization or automation—achieves comparable results.

Self-refinement in LLMs. Our work connects to the broader literature on LLM self-improvement. Self-Refine Madaan et al. [2023] demonstrates that LLMs can iteratively improve their own outputs through self-generated feedback, achieving roughly 20% preference gains. Our iterative rejection protocol is structurally similar—generate, receive feedback, regenerate—but the feedback is external (a human saying “too AI”) rather than self-generated, and the optimization target is detector evasion rather than general quality.

Positioning our work. table 1 summarizes how our approach differs from prior evasion methods. We are the first to study the simplest possible attack: a non-technical user iteratively rejecting outputs through normal conversation. Our best-of- N control further distinguishes this from prior work by isolating contextual steering from selection pressure.

3 Methodology

We describe the dataset, experimental conditions, detection pipeline, and statistical analysis plan.

3.1 Dataset

We sample 50 questions from the `open_qa` split of the HC3 (Human ChatGPT Comparison Corpus) Guo et al. [2023] using a fixed random seed (42). HC3 provides human-written answers for each question, which serve as calibration baselines for the detectors. We chose open-domain question answering because it represents a common use case for LLM-generated text and is well-represented in prior detection benchmarks Dugan et al. [2024].

3.2 Generation Conditions

Each question is answered under five conditions using GPT-4.1 (`gpt-4.1`, temperature = 1.0, seed = 42):

Vanilla. A standard system prompt with no evasion instructions. This establishes the upper bound on detection scores and serves as the primary baseline.

Single-step humanization. The system prompt instructs the model to “write naturally and avoid AI patterns.” This represents a naive, zero-effort evasion attempt.

Strong single-step. A detailed system prompt listing specific AI markers to avoid (e.g., “do not use words like *delve*, *crucial*, *landscape*”; “use contractions”; “vary sentence length”). This tests whether a well-crafted single prompt can match iterative rejection.

Iterative rejection. The model generates an initial response, then receives five rounds of escalating rejection feedback with full conversation history preserved:

1. “This sounds too AI-generated. Rewrite it to sound more like a real person wrote it.”
2. “Still sounds like AI. Try again—write like an actual human would.”
3. “Nope, still reads like ChatGPT output. Make it sound genuinely human.”
4. “This still has that AI quality to it. One more try—really make it sound natural.”
5. “Getting closer but still detectably AI. Give it one final shot, as human-sounding as possible.”

This yields six outputs per question (initial + five revisions), all scored by both detectors.

Best-of-5 selection. Five independent vanilla generations per question. We select the output with the lowest BINOCULARS score. This control isolates selection pressure: if iterative rejection merely increases the chance of sampling a low-scoring output, best-of-5 should perform comparably.

In total, we make approximately 750 API calls to GPT-4.1.

3.3 Detection Pipeline

We score all generated text with two detectors that represent different detection paradigms:

ROBERTA classifier. The `openai-community/roberta-base-openai-detector` model So-laiman et al. [2019], a fine-tuned RoBERTa-base classifier trained on GPT-2 era text. It outputs $P(\text{AI}) \in [0, 1]$. We include it as a representative trained classifier, while noting its known limitations on modern LLM outputs.

BINOCULARS zero-shot detector. The method of Hans et al. [2024], which computes a perplexity ratio between Falcon-7B and Falcon-7B-Instruct. We use the accuracy-optimized threshold of 0.9015 and convert raw scores to $P(\text{AI})$ via sigmoid mapping. BINOCULARS represents the current state of the art in zero-shot detection and is our primary evaluation metric.

Both detectors run on $4 \times$ NVIDIA RTX A6000 GPUs (49 GB each) using PyTorch 2.10.0.

3.4 Evaluation Metrics

Primary metric. Mean $P(\text{AI})$ across 50 questions per condition (continuous score). Lower values indicate more successful evasion.

Statistical tests. We use paired Wilcoxon signed-rank tests, which are non-parametric and appropriate for paired, non-normally distributed detection scores. We apply Bonferroni correction for four pairwise comparisons ($\alpha = 0.0125$). We report matched-pairs rank-biserial correlation (r) as the effect size measure.

Linguistic features. To understand the mechanism of evasion, we track five linguistic features across rejection rounds: type-token ratio (vocabulary diversity), first-person pronoun rate, contraction rate, AI marker word count (e.g., “delve,” “crucial,” “furthermore”), and average sentence length.

4 Results

We present detection scores across conditions (section 4.1), the per-round trajectory of iterative rejection (section 4.2), the comparison with best-of-5 selection (section 4.3), and linguistic feature analysis (section 4.4).

4.1 Main Detection Results

table 2 and table 3 show mean detection scores across all conditions for both detectors.

Three findings stand out. First, iterative rejection substantially reduces detection: from 0.798 to 0.483 on BINOCULARS (−39.5%). Second, the strong single-step prompt achieves an even larger reduction to 0.342 (−57.1%), approaching human text levels. Third, the naive single-step prompt (i.e., “write naturally”) is counterproductive, slightly *increasing* detection scores on BINOCULARS (0.835 vs. 0.798).

Figure 1 shows the distribution of scores across conditions, illustrating that the strong single-step and iterative rejection conditions shift the score distribution substantially leftward.

Table 2: Detection scores by condition (BINOCULARS, zero-shot). Lower $P(\text{AI})$ indicates more successful evasion. Best AI-generated result in **bold**. Human text shown for calibration.

Condition	Mean $P(\text{AI})$	Std	Δ vs. Vanilla
Human text	0.325	0.376	—
Vanilla	0.798	0.236	baseline
Single-step	0.835	0.228	+0.037
Strong single-step	0.342	0.292	−0.456
Iterative (R5)	0.483	0.280	−0.315
Best-of-5 (selected)	0.702	0.286	−0.096

Table 3: Detection scores by condition (ROBERTA, trained classifier). This detector assigns human text a *higher* AI probability (0.320) than vanilla GPT-4.1 (0.264), indicating poor generalization to modern LLM outputs. Results included for completeness but should be interpreted with caution.

Condition	Mean $P(\text{AI})$	Std	Δ vs. Vanilla
Human text	0.320	0.334	—
Vanilla	0.264	0.387	baseline
Single-step	0.294	0.400	+0.030
Strong single-step	0.039	0.147	−0.225
Iterative (R5)	0.082	0.226	−0.182
Best-of-5 (selected)	0.103	0.257	−0.161

4.2 Iteration Trajectory

Figure 2 shows the mean BINOCULARS score at each rejection round. The trajectory is non-monotonic, with the largest single drop occurring between rounds 2 and 3 (−0.203), followed by a partial rebound at round 4.

table 4 reports the per-round scores. The overall trend is downward (R0: 0.855 → R5: 0.483), but the path is uneven. The R2→R3 transition accounts for more than half of the total improvement, suggesting a tipping point in how the model responds to accumulated rejection context.

4.3 Best-of-5 vs. Iterative Rejection

A central question is whether iterative rejection provides genuine steering or merely benefits from selection pressure (i.e., generating multiple outputs increases the chance of a low-scoring one). We control for this with the best-of-5 condition, which generates five independent vanilla outputs and selects the lowest-scoring one.

table 5 reports the statistical comparisons. Iterative rejection significantly outperforms best-of-5 selection on BINOCULARS (0.483 vs. 0.702, $p < 0.001$, $r = 0.555$), confirming that the rejection context provides steering beyond selection. The iterative approach achieves lower scores than best-of-5 on 36 of 50 questions (72%).

Figure 3 shows a per-question scatter plot comparing iterative rejection against best-of-5 selection. Most points fall below the diagonal, indicating iterative rejection achieves lower scores on the majority of questions.

4.4 Linguistic Feature Analysis

To understand *how* the model changes its generation across rejection rounds, we track five linguistic features. table 6 reports the values at the initial round (R0), final round (R5), and for human text.

Figure 4 visualizes the per-round trajectories. Two patterns are noteworthy. First, the type-token ratio increases monotonically toward human levels, indicating richer vocabulary. Second, first-person pronoun usage increases 26× by R5, dramatically overshooting the human baseline (0.016




Figure 1: Distribution of BINOCULARS detection scores across conditions. The strong single-step prompt and iterative rejection (R5) both shift the distribution toward human-level scores, while the naive single-step prompt remains close to vanilla.

Table 4: Per-round BINOCULARS scores during iterative rejection. The R2→R3 transition shows the largest single-round improvement.

Round	Mean $P(\text{AI})$	Δ from R0
R0 (initial)	0.855	—
R1	0.722	−0.133
R2	0.724	−0.131
R3	0.521	−0.334
R4	0.600	−0.255
R5 (final)	0.483	−0.372

vs. 0.002). The model does not simply converge toward human statistics; it enters an exaggerated “casual first-person” mode.

Critically, the features the model *fails* to change are equally informative. AI marker words (“delve,” “crucial,” “landscape”) remain at 0.24 per text, twelve times the human rate. Contraction rates actually decrease, contradicting the intuition that contractions sound more human. These failures suggest the model does not have fine-grained control over individual stylistic markers through rejection feedback alone.

5 Discussion

5.1 Regime Shift vs. Feature Suppression

Our linguistic analysis supports a *regime shift* interpretation over a *feature suppression* interpretation of iterative rejection.

What would feature suppression look like? If the model were learning to suppress specific detectable features through rejection feedback, we would expect AI marker words to decrease, contractions to increase, and the overall linguistic profile to converge toward human norms monotonically. None of these occur.




Figure 2: Mean BINOCULARS $P(\text{AI})$ across rejection rounds. The sharpest improvement occurs between R2 and R3, suggesting a phase transition rather than gradual convergence. The non-monotonic pattern ($R4 > R3$) indicates the model oscillates between generation regimes.

Table 5: Pairwise statistical comparisons against iterative rejection (R5) on BINOCULARS scores. All comparisons are significant after Bonferroni correction ($\alpha = 0.0125$). Effect sizes (r) are matched-pairs rank-biserial correlations.

Comparison	p -value	Effect size (r)	Significant
Single-step vs. Iterative	1.08×10^{-9}	0.879	Yes
Strong single-step vs. Iterative	8.09×10^{-3}	0.426	Yes
Best-of-5 vs. Iterative	4.51×10^{-4}	0.555	Yes
Vanilla vs. Iterative	2.47×10^{-7}	0.776	Yes

What do we observe instead? The model enters a qualitatively different generation mode. First-person pronoun usage increases $26\times$, far exceeding human norms. The model adds conversational meta-commentary (“Alright, here’s my most natural shot:”) and adopts an informal, anecdotal style. Meanwhile, AI marker words remain unchanged and contractions actually decrease. The non-monotonic trajectory—with a sharp drop at R3 and a partial rebound at R4—further suggests the model oscillates between regimes rather than converging smoothly.

This distinction matters for detector design. If evasion worked through targeted feature suppression, detectors could simply weight those features more heavily. A regime shift, however, produces a qualitatively different distribution that requires detectors to model the “trying to sound human” mode as a distinct class.

5.2 Why Generic Humanization Fails

A surprising finding is that the naive single-step prompt (“write naturally, avoid AI patterns”) slightly *increases* BINOCULARS detection scores (0.835 vs. 0.798 vanilla). We hypothesize that vague humanization instructions trigger a “trying too hard” mode: the model adds filler, qualifications, and an unnatural self-consciousness that is more detectable, not less.

In contrast, the strong single-step prompt succeeds because it provides *specific, actionable* instructions (avoid particular words, use contractions, vary sentence length). This finding has practical implications: for users seeking to evade detectors, prompt specificity matters more than iteration




Figure 3: Per-question comparison of BINOCULARS scores: iterative rejection (R5) vs. best-of-5 selection. Points below the diagonal indicate questions where iterative rejection achieves lower (better) scores. The iterative approach wins on 72% of questions.

Table 6: Linguistic features across rejection rounds (BINOCULARS). The model increases vocabulary diversity and first-person pronoun usage (overshooting human norms) while failing to reduce AI marker words or increase contractions.

Feature	R0	R5	Human	Trend
Type-token ratio	0.668	0.749	0.826	↑ toward human
First-person rate	0.0006	0.016	0.002	↑↑ overshoots
Contraction rate	0.004	0.0004	0.005	↓ away from human
AI marker count	0.24	0.24	0.02	→ no change
Avg. sentence length	20.0	19.3	21.8	→ minimal change

count. For detector developers, this suggests that prompt-engineering-based evasion may be the more pressing threat.

5.3 Implications

For detector trustworthiness. Both iterative rejection and strong single-step prompting reduce detection to near-human levels. The strong single-step approach is particularly concerning because it requires zero iteration—a single well-crafted system prompt achieves 0.342 on BINOCULARS, within 0.017 of the human baseline (0.325). Detector-based enforcement alone is insufficient for high-stakes decisions.

For alignment research. The regime shift phenomenon suggests that LLMs harbor distinct generation modes activatable through conversational context. The model does not simply adjust feature weights; it switches to a different “persona.” This connects to work on jailbreaking and persona-based steering Madaan et al. [2023] and suggests that conversational context exerts a qualitative, not just quantitative, influence on generation.




Figure 4: Linguistic feature trajectories across rejection rounds. Type-token ratio converges toward human levels, while first-person pronoun usage overshoots dramatically. AI marker words and contraction rates remain largely unchanged, indicating a regime shift rather than targeted feature suppression.

For the ROBERTA detector. The ROBERTA classifier assigns human text a higher AI probability (0.320) than vanilla GPT-4.1 output (0.264). This detector, trained on GPT-2 era text, has effectively negative discriminative power for modern LLMs. We recommend that practitioners discontinue its use for any text potentially generated by GPT-4-class models.

5.4 Limitations

Single LLM. We test only GPT-4.1. The regime shift phenomenon may differ across model families (Claude, Gemini, open-source models like LLaMA-3). Replication across models is needed.

Sample size. Our 50-question dataset is sufficient for paired statistical tests but limits power for subgroup analyses (e.g., by question type or answer length).

Fixed rejection protocol. We use the same five escalating rejection messages for all questions. Varying the specificity of rejection feedback (e.g., “you used ‘delve’ and ‘furthermore’ ” vs. generic “too AI”) could reveal which aspects of feedback are most effective.

Detector coverage. We test two open-source detectors. Commercial systems (GPTZero, Turnitin) may respond differently. The BINOCULARS sigmoid calibration may affect absolute $P(\text{AI})$ values, though relative comparisons remain valid.

Temperature and sampling. All generation uses temperature = 1.0 with seed = 42. Different temperature settings could change both the baseline detection rate and the effectiveness of iterative rejection.

6 Conclusion

We studied the simplest possible evasion attack against AI text detectors: iteratively telling an LLM “this sounds too AI-generated, try again.” Using GPT-4.1 on 50 open-domain questions scored by two detectors, we find that:

1. **Iterative rejection works.** Five rounds of conversational rejection reduce BINOCULARS detection scores by 43% ($p < 0.001$).

2. **It provides genuine steering, not just selection.** Iterative rejection significantly outperforms best-of-5 selection ($p < 0.001$, $r = 0.56$), confirming that the rejection context changes how the model generates, not just how many chances it gets.
3. **A well-crafted single prompt works even better.** A strong single-step prompt with specific anti-AI instructions achieves a 60% reduction, approaching human text levels—suggesting that prompt specificity matters more than iteration.
4. **The mechanism is a regime shift.** The model does not suppress specific AI features; it switches to a distinct “trying to sound human” persona characterized by excessive first-person narration and higher vocabulary diversity, while AI marker words persist.

These findings have immediate practical implications: detector-based enforcement is fragile against trivial attacks that require no technical expertise. For detector developers, robustness to persona shifts—not just vocabulary-level features—is essential. For alignment researchers, the regime shift phenomenon reveals that conversational context can qualitatively change LLM generation in ways that simple feature-level analysis may miss.

Future work should replicate these findings across model families, test commercial detectors, vary the specificity of rejection feedback, and use probing analysis to directly measure internal regime shifts. The arms race between detection and evasion continues, but the barrier to effective evasion is lower than previously assumed.

References

- Guangsheng Bao, Yanbin Zhao, Zhiyang Teng, Linyi Yang, and Yue Zhang. Fast-DetectGPT: Efficient zero-shot detection of machine-generated text via conditional probability curvature. In *International Conference on Learning Representations (ICLR)*, 2024.
- Zhongtao Cheng, Vinu Sankar Sadasivan, et al. Adversarial paraphrasing: A universal attack for humanizing AI-generated text. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- Liam Dugan, Alyssa Hwang, Filip Trhlík, Andrew Lasko, Ondřej Stamm, Minkee Ip, Jack Zhu, Nikita Juri, and Chris Callison-Burch. RAID: A shared benchmark for robust evaluation of machine-generated text detectors. In *Association for Computational Linguistics (ACL)*, 2024.
- Yifan Fang et al. CoPA: Contrastive paraphrase attacks on LLM-generated text detectors. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2025.
- Biyang Guo, Xin Zhang, Ziyuan Wang, Minqi Jiang, Jinran Nie, Yuxuan Ding, Jianwei Yue, and Yupeng Wu. HC3: Human ChatGPT comparison corpus. *arXiv preprint arXiv:2301.07597*, 2023.
- Abhimanyu Hans, Avi Schwarzschild, Valeriia Cheber, Wojciech Czaja, Micah Goldblum, and Tom Goldstein. Binoculars: Zero-shot detection of LLM-generated text. In *International Conference on Machine Learning (ICML)*, 2024.
- John Kirchenbauer, Jonas Geiping, Yuxin Wen, Jonathan Katz, Ian Miers, and Tom Goldstein. A watermark for large language models. In *International Conference on Machine Learning (ICML)*, 2023.
- Kalpesh Krishna, Yixiao Song, Marzena Karpinska, John Wieting, and Mohit Iyyer. Paraphrasing evades detectors of AI-generated text, but retrieval is an effective defense. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Colin Lu et al. SICO: Large language models can be guided to evade AI-generated text detection. *Transactions on Machine Learning Research (TMLR)*, 2024.
- Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. Self-refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D Manning, and Chelsea Finn. DetectGPT: Zero-shot machine-generated text detection using probability curvature. In *International Conference on Machine Learning (ICML)*, 2023.

- Irene Solaiman, Miles Brundage, Jack Clark, Amanda Aspell, Ariel Herbert-Voss, Jeff Wu, Alec Radford, Gretchen Krueger, Jong Wook Kim, Sarah Kreps, et al. Release strategies and the social impacts of language models. *arXiv preprint arXiv:1908.09203*, 2019.
- Ying Zhou, Ben He, and Le Sun. Humanizing machine-generated content: Evading AI-text detection through adversarial attack. *arXiv preprint arXiv:2404.01907*, 2024.
- Yiran Zhou et al. Self-disguise attack: Evading AI-generated text detection via adversarial feature extraction. *arXiv preprint arXiv:2508.15848*, 2025.